



Afraz Khan

AI ENGINEER

afrazkhan1407@gmail.com +92 348 646 5743 Islamabad, Pakistan LinkedIn GitHub Upwork

SUMMARY

AI Engineer based in Islamabad, specializing in LLMs, multi-agent systems, and computer vision. I design and deploy scalable ML solutions — from multi-agent LLM pipelines and semantic search to OCR-powered document automation — for healthcare and FinTech. My work has slashed operational costs by €40,000, boosted query accuracy by 40%, and cut customer support response times by 75%.

EXPERIENCE

AI Engineer

Present

OphyCare · Islamabad, Pakistan (Hybrid · Full-time)

- Built an end-to-end medical scribing pipeline converting patient-doctor interactions into structured clinical data.
- Automated population of 10+ mandatory medical forms with context-aware Generative AI, eliminating hours of manual entry per provider.
- Deployed multi-agent systems where specialized agents collaborate to resolve complex healthcare workflows including eligibility verification.
- Integrated advanced GenAI solutions to streamline the most time-consuming parts of clinical documentation.

Python LangGraph OpenAI Deepgram pgvector FastAPI

AI / ML Engineer

Apr 2025 — May 2025

EDGE C Pvt. Ltd · Lahore, Pakistan (Hybrid · Full-time)

- Drove a 40% uplift in query accuracy over legacy approaches using Neo4j knowledge graphs and MongoDB vector embeddings.
- Built LangChain orchestration pipelines delivering a 35% improvement in answer relevance for complex financial queries.
- Developed gRPC microservices with Redis caching achieving sub-200 ms end-to-end retrieval latency.
- Collaborated with data engineering and product teams to define financial ontology schemas and optimize embedding strategies.

Python LangChain Neo4j MongoDB Redis gRPC

AI Engineer

Sep 2024 — Jan 2025

Ace Money Transfer · Kharian, Pakistan (Full-time)

- Designed and implemented ML models for critical financial operations and transaction analytics.
- Engineered robust data preprocessing pipelines for complex, high-volume transactional datasets.

Python Scikit-learn PyTorch Data Pipelines

Freelance AI Engineer

Ongoing

Upwork & Fiverr · Remote

- Delivered custom AI/ML solutions for clients globally — from rapid prototypes to production-grade systems.
- Projects span LLM applications, RAG pipelines, chatbot development, OCR pipelines, and intelligent automation.

LLM Systems RAG Computer Vision OCR NLP APIs

AI Intern

Prior

Rev9 Solutions · Jhelum, Pakistan (Internship)

- Engineered ML models to address client business challenges across industries under senior AI engineers.
- Transformed raw data into structured formats for model training and validation.
- Integrated solutions into production environments alongside senior AI engineers.

Python Scikit-learn TensorFlow

EDUCATION

BS, Computer Science

University of Punjab · Lahore, Pakistan

CGPA 3.8 / 4.0

SKILLS

LLMs & Agents	OpenAI, LangGraph, LangChain, Multi-Agent Systems, RAG, Semantic Search, Prompt Engineering
Machine Learning	PyTorch, TensorFlow, Keras, Scikit-learn, Deep Learning, NLP, BERT / Transformers, Hugging Face
Computer Vision	PaddleOCR, Tesseract, YOLOv8, OpenCV, DeepFace, Image Processing, Document Analysis
Backend & APIs	Python, FastAPI, Flask, gRPC, REST APIs, Docker, Microservices
Data & Infrastructure	Neo4j, pgvector, MongoDB, Redis, PostgreSQL, Vector DBs, Knowledge Graphs
MLOps & Tools	Git, Deepgram, Roboflow, LabelImg, Model Optimization, Data Visualization

SELECTED PROJECTS

Eligibility Verification Agent Automated insurance eligibility for healthcare - AI agent automating the full eligibility pipeline — clearinghouse inquiry → EMR updates → physician sign-off. - Eliminates manual back-and-forth between clinical staff, clearinghouses, and EMR systems.	Mar 2026 — Present
Sadiq Scribe AI medical scribing & form automation - Transcribes patient-doctor audio and auto-populates 10+ mandatory medical forms. - End-to-end pipeline: real-time STT → clinical entity extraction → intelligent form-field mapping.	Jul 2025 — Feb 2026
SMR — Smart Medical Record Clinical Q&A with citation support - Hybrid retrieval combining semantic search with structured lookups surfaces citation-backed answers. - Physicians query current and historical appointment data with proper source attribution.	Oct 2025 — Feb 2026
Intellectus Hybrid knowledge retrieval for FinTech - Dual-module system combining graph-based and semantic search — 40% uplift in query accuracy. - Neo4j for entity relationships, MongoDB vectors for unstructured content, LangChain orchestration (+35% answer relevance).	Apr 2025 — May 2025
KYC — Know Your Customer End-to-end verification system for FinTech - Slashed annual operational costs by €40,000 with a full-stack KYC solution. - Integrated PaddleOCR and DeepFace for document verification and facial recognition.	Sep 2024 — Jan 2025
Doc Detector & Classifier Lightweight document analysis for FinTech 98% field detection accuracy on identity documents; trained on 5,000+ custom-labeled documents. - Lightweight YOLOv8n model reduced compute requirements by 70% for edge deployment.	Oct 2024 — Nov 2024
Knowledge Base Q&A Intelligent information retrieval engine - NLP-powered Q&A system that cut customer support response times by 75%. - Hybrid TF-IDF + BERT retrieval; Flask API handling 1,000+ concurrent queries at sub-100ms.	Jan 2024 — Feb 2024